

MaskClaw: Edge-Side Personalized Privacy Arbitration for GUI Agents with Behavior-Driven Skill Evolution

Yanqiu Zhao^{1,†} Dongying Zheng^{1,†} Kaibo Huang^{1,*} Yukun Wei¹ Zhongliang Yang^{1,*} Linna Zhou¹

¹Beijing University of Posts and Telecommunications

[†]Equal contribution. Authors are listed in alphabetical order. ^{*}Corresponding authors.
{ZYQiu, zhengdongying, Huangkaibo, weiyukun, yangzl, zhoulinna}@bupt.edu.cn

Abstract

GUI agents rely on screenshots to infer intent and operate across applications, but these screenshots often contain private messages, medical records, payment credentials, and workplace-specific workflows. Privacy decisions in this setting depend on task, recipient, application state, and user role, yet static PII detectors miss these boundaries and cloud-side VLM reasoning can upload the raw screen before deciding what should be protected. We present MaskClaw, an edge-side privacy arbitrator for GUI agents. MaskClaw extracts local visual evidence, retrieves user- and task-specific policy memory, and decides Allow, Mask, or Ask before raw screenshots leave a trusted user- or organization-controlled environment. In five designed skill-evolution scenarios, it turns corrections, cancellations, and edits into reusable privacy skills checked by a sandbox gate. We introduce P-GUI-Evo, a benchmark built from real UI patterns, reconstructed HTML screens, and sanitized labels. Experiments show that pattern matching, cloud reasoning, and routing alone tend to over-confirm, over-mask, or expose raw screenshots under the same protocol. The artifact is available at <https://github.com/Theodora-Y/MaskClaw..>

1 Introduction

Graphical user interface (GUI) agents automate mobile and desktop tasks by interpreting screenshots and acting on user instructions (Zhang et al., 2025; Liu et al., 2024; Wang et al., 2024). Multimodal foundation models and GUI grounding methods drive this progress (Hong et al., 2024; Zheng et al., 2024; Wang et al., 2025b), while compact multimodal models make edge-side arbitration feasible for extracting local evidence and deciding what observation may leave the device (Hu et al., 2024; Yao et al., 2024). The edge side may be the endpoint itself or a trusted local environment where the AI




Figure 1: Overview of the motivation and problem setting. GUI agents need screenshots to act, but the same screenshots can expose task-irrelevant private content. MaskClaw treats privacy as edge-side arbitration: decide Allow, Mask, or Ask before raw screenshots leave the trusted environment.

model, large-language-model component, or privacy computation runs, including phones, PCs, in-vehicle systems, smart-home devices, workstations, personal servers, and private enterprise servers.

The same visual interface creates a privacy bottleneck. Screenshots may contain private messages, account identifiers, medical records, financial tables, business procedures, or task-irrelevant personal details. Prior work shows that language and vision-language models can leak, infer, or be induced to reveal sensitive information from textual and visual inputs (Carlini et al., 2021; Kim et al., 2023; Tömekçe et al., 2024; Wang et al., 2026b; Zhang et al., 2026). Sending raw screenshots to a remote VLM creates a cloud-dependency paradox: the screen is exposed before protection is decided.

Existing privacy mechanisms leave this decision boundary unresolved. Pattern-based tools detect common PII (Habernal et al., 2023), and prompt-side masking can reduce online LLM exposure (Zhu et al., 2024). GUI-agent privacy depends on context and personal policy. A value that is routine in a reimbursement form may require protection when it appears as a bank balance, and the same screen can be released for one persona and re-

quire confirmation for another. The correct action depends on screen content, task intent, recipient, application state, and user role (Lu et al., 2024b; Chaudhari et al., 2026; Wang et al., 2025a; Salemi et al., 2023). We arbitrate exposure on the trusted edge before releasing raw screens, with policies refined from user behavior.

In summary, our main contributions are:

- We propose MaskClaw, an edge-side privacy arbitrator that uses local evidence and policy memory to choose Allow, Mask, or Ask before raw screenshots leave a trusted environment.
- We introduce a feedback-to-skill loop that turns corrections, cancellations, edits, and negative feedback into reusable privacy skills checked by rule tests, held-out cases, and sandbox validation.
- We construct P-GUI-Evo, a sanitized, HTML-reconstructed benchmark for multi-persona GUI privacy, covering local policy memory, rule-free routes, cloud and routing controls, and skill evolution under one protocol.

2 Related Work

2.1 GUI grounding for agentic screen control

GUI agents such as AppAgent, AutoGLM, Mobile-Agent, CogAgent, and Step-GUI combine visual state understanding with action execution (Zhang et al., 2025; Liu et al., 2024; Wang et al., 2024; Hong et al., 2024; Yan et al., 2025). GUI grounding methods and benchmarks structure screenshots into actionable observations and evaluate device-control behavior (Yang et al., 2023; Lu et al., 2024a; You et al., 2024; Zheng et al., 2024; Rawles et al., 2023; Xu et al., 2025). MaskClaw builds on these perception and control capabilities to govern which visual evidence may be released to a downstream or cloud agent.

2.2 Contextual privacy for visual agent observations

Closest to our setting, VLM and GUI-agent privacy benchmarks show that visual inputs can leak private or socially contextual information (Lu et al., 2024b; Chaudhari et al., 2026; Tömekçe et al., 2024; Wang et al., 2026b; Zhang et al., 2026; Wang et al., 2026a). GUI anonymization uses type-preserving

placeholders to keep interfaces usable while hiding visual content (Zhao et al., 2026); MemPrivacy protects edge-cloud agent memory through local privacy-span extraction and type-aware placeholder substitution (Chen et al., 2026). Agentic settings add further risk because prompt injection and privacy jailbreaks can turn one exposed observation into downstream leakage (Chang et al., 2026; Wang et al., 2025c). These methods reduce exposure while leaving the live Allow/Mask/Ask decision open under a specific task, recipient, application state, and user role. MaskClaw treats this choice as context-conditioned policy arbitration, consistent with contextual integrity (Nissenbaum, 2004), before raw screenshots leave the trusted edge boundary.

2.3 Feedback-driven optimization of agent policies

Language feedback can optimize prompts, reasoning traces, and agent behavior through textual search, as shown by OPRO, TextGrad, Reflexion, STaR, evolutionary prompt optimization, and WebRL (Yang et al., 2024; Yuksekogonul et al., 2024; Shinn et al., 2023; Zelikman et al., 2022; Guo et al., 2024; Qi et al., 2025). Recent GUI skill systems explore verifiable tree evolution, recursive skill-RL, collective skill evolution, and memory-designed reusable skills (Jiang et al., 2026; Xia et al., 2026; Ma et al., 2026; Zhou et al., 2026). MaskClaw re-purposes this paradigm for personalized privacy skills, keeping candidate updates only after sandbox validation confirms policy alignment and safe exposure control.

MaskClaw connects GUI grounding, privacy protection, and skill optimization through local evidence, context-conditioned arbitration, and auditable privacy-skill updates.

3 Problem Formulation and Benchmark

3.1 Personalized Privacy Decision

We formulate GUI privacy arbitration as a personalized, context-conditioned classification problem. The arbitrator decides whether visual evidence may be exposed under a specific user, task, recipient, and application context. The same screen element can receive different decisions across personas or workflows because privacy depends on user-specific context beyond input content alone (Salemi et al., 2023; Zhang et al., 2026).

For the i -th GUI-agent interaction, we define the

prediction instance as:

$$x_i = (I_i, q_i, a_i, c_i, p_i), \quad (1)$$

where I_i is the current GUI screenshot, q_i the user instruction or task goal, a_i the candidate intent or next action of the GUI agent, c_i the application platform and relational context, and p_i the persona and policy context. An arbitrator predicts $y_i = f(x_i)$ with output space $\mathcal{Y} = \{\text{Allow}, \text{Mask}, \text{Ask}\}$, corresponding to no intervention, required protection, and required user consent.

3.2 The P-GUI-Evo Benchmark

To evaluate personalized privacy arbitration and behavior-driven skill evolution, we construct P-GUI-Evo, a real-UI-pattern-derived, HTML-reconstructed benchmark of 832 GUI privacy samples. The benchmark covers 296 normalized scenario cores, with 2.81 samples per core on average. P-GUI-Evo connects GUI-agent and mobile-control benchmarks, which focus on interaction and task completion, with VLM and GUI-agent privacy benchmarks, which focus on visual leakage or privacy assessment (Rawles et al., 2023; Xu et al., 2025; Lu et al., 2024b; Wang et al., 2026b; Zhang et al., 2026; Wang et al., 2026a). It evaluates Allow/Mask/Ask decisions from GUI observations and task context, including cases where similar visual evidence requires different decisions across personas or workflows.

We build each record through a controlled reconstruction and audit workflow. We start from common GUI interaction patterns and platform-like layouts, using placeholders, fake identities, or empty form states as sanitized surrogates for private content. An LLM drafts the HTML structure and scenario text and supports PII-category consistency checks. The project authors fix task content and audit plausibility, fluency, alignment with the candidate agent intent, and support for the expected decision. OCR checks and PII audits verify that visible sensitive evidence can be detected and that the rendered screenshot matches the scenario metadata.

Each sample contains a GUI screenshot, application platform, persona context, relationship context, user instruction, agent intent, benchmark bucket, and an expected Allow/Mask/Ask decision for offline evaluation. The final joint-label distribution is Mask=438, Allow=314, and Ask=80. Prediction methods receive only runtime-available fields; risk

Partition	Count	Purpose
User A	234	healthcare and medical workflow
User B	246	commerce and livestream selling
User C	352	office and everyday services
D1 basic	174	clean in-distribution cases
D2 generalization	546	UI/task variants
D3 stress	112	noisy boundary cases

Table 1: Persona and bucket distribution in P-GUI-Evo.

Construction or audit statistic	Count
Active benchmark samples	832
Normalized scenario cores	296
Cores with expanded variants	144
Samples with non-None PII type	715
Location-ready PII evidence samples	685
Localizable evidence boxes	1,529
L1/L3 usable samples	736

Table 2: Construction and audit statistics for P-GUI-Evo. PII types are sample-level semantic labels; evidence boxes are location-level visual items.

tags, PII categories, and lower-level audit labels are reserved for analysis. Scenario cores are expanded into structural, visual, or lexical variants that preserve the privacy boundary. Among the 296 cores, 144 have at least one expanded variant, with 173 lexical, 201 structural, and 208 visual-shift rows.

Table 1 summarizes the persona and bucket distribution. User A/B/C are synthetic persona identifiers for varying workflow, recipient, and task context. D3 stress cases include OCR noise, occlusion, low resolution, mixed languages, popup interference, or multiple sensitive regions.

Table 2 summarizes construction-audit coverage. All 832 samples include complete task, context, decision, masking, and PII-type metadata. Sample-level PII annotations cover medical, address, financial, identity, contact, credential, compensation, supplier/customer, and security-related signals; the largest localizable evidence categories are medical, financial, contact, address/location, and account or identity credentials.

4 The MaskClaw Framework

MaskClaw is an edge-side privacy and personalization layer between the live GUI and downstream agents. It arbitrates each outgoing observation inside a trusted user- or organization-controlled environment and refines auditable scenario-level privacy skills from validated feedback.

4.1 Framework Overview

MaskClaw has two coupled loops. In the online loop, the edge-side controller keeps raw screenshots local, extracts evidence, selects an exposure action through policy arbitration, and controls what leaves the device through SafeScreenshot construction. The evolution loop converts confirmations, cancellations, edits, and rejected actions into candidate skill updates that enter local memory after validation. Execution semantics remain fixed, and inspectable policy memory supports personalization.

Figure 2 summarizes the pipeline. We next describe its four components.

4.2 Local Perception and Evidence Extraction

Perception extracts screen evidence for privacy reasoning before any raw screenshot leaves the device. It produces structured evidence for arbitration and SafeScreenshot construction.

For the current interaction, denoted by s , MaskClaw extracts an evidence set:

$$E(s) = \{e_1, e_2, \dots, e_n\}. \quad (2)$$

Each item may contain OCR text, field type, semantic description, confidence, and, when available, a bounding box. Localized evidence supports redaction. Non-localized evidence informs arbitration, following GUI parsing and grounding work (Yang et al., 2023; Zheng et al., 2024). Perception proposes evidence, and arbitration decides exposure.

4.3 Rule-Grounded Policy Arbitration

Arbitration is the fixed policy interface between local perception and the exposure action. MaskClaw grounds this interface in local policy memory:

$$C_t = \{r_1, r_2, \dots, r_m\}, \quad (3)$$

where each entry is represented at the interface level as

$$r_j = (s_j, t_j, a_j, p_j, z_j). \quad (4)$$

Here s_j is the rule scope, covering application, persona or role, recipient boundary, and intended action; t_j is the trigger evidence, such as OCR text, field type, semantic description, or a localizable region; a_j is the recommended exposure action; p_j is a conflict-resolution priority; and z_j is a short rationale. This schema keeps policy conditions inspectable at the interface level.

For each interaction, MaskClaw builds a retrieval query from application, persona, recipient

boundary, intended action, and local visual evidence $E(s)$. It retrieves a small candidate set $R_k(s)$ from C_t using retrieval-augmented generation with local, user-specific memory (Lewis et al., 2020; Asai et al., 2024). Candidate rules are ranked by facet match and evidence compatibility, with recipient, action, and persona matches weighted above generic field-type matches.

The controlled arbitrator selects an applicable entry r^* from $R_k(s)$ and emits a masking variable $m(s)$ and a confirmation variable $u(s)$ as locked fields. More specific persona-recipient-action rules override generic rules. When the retrieved set lacks a sufficiently applicable rule, the arbitrator falls back to conservative boundary tests over task necessity, recipient trust, and sensitive evidence. The final policy mapping is deterministic:

$$\text{Decision}(s) = \begin{cases} \text{Ask}, & u(s) = 1, \\ \text{Mask}, & u(s) = 0 \wedge m(s) = 1, \\ \text{Allow}, & \text{otherwise.} \end{cases} \quad (5)$$

Confirmation has priority over masking, and masking has priority over allowing. The Allow/Mask/Ask interface remains fixed and auditable as policy memory evolves.

4.4 SafeScreenshot Construction

SafeScreenshot construction turns the arbitration result into an observation boundary. Let $B(s)$ be the localizable regions from perception. Under Mask, MaskClaw uses the selected policy rule and local evidence to choose protected regions:

$$B^*(s) \subseteq B(s). \quad (6)$$

It then constructs a mediated screenshot:

$$I_{\text{safe}} = \text{redact}(I, B^*(s)). \quad (7)$$

The redaction operation can blur, mask, replace, or otherwise transform protected regions locally. Policy gating lets the same field remain visible, be masked, or trigger confirmation depending on the task, recipient, and user context. The mediated screenshot preserves task-relevant layout and non-sensitive content and weakens protected content.

Under Ask, MaskClaw pauses execution and returns a confirmation request with its rationale. Under Allow, the task proceeds through the controlled observation route with raw-screen access mediated by arbitration.




Figure 2: MaskClaw architecture. The edge-side layer keeps raw screenshots local, extracts visual evidence, retrieves persona- and task-conditioned policy rules, and arbitrates an Allow/Mask/Ask exposure action. After arbitration, downstream agents receive allowed content, SafeScreenshot outputs, or confirmation requests. Validated user feedback updates auditable privacy skills.

4.5 Behavior-Driven Skill Evolution

Skill evolution converts post-decision feedback into persistent local memory (Packer et al., 2023). MaskClaw records traces:

$$\tau = [(s_1, d_1, y_1), \dots, (s_T, d_T, y_T)], \quad (8)$$

where s_t is the GUI state, d_t the system decision, and y_t contains downstream actions and feedback such as confirmations, cancellations, edits, rejected actions, retries, or explicit skill instructions.

From these traces, MaskClaw extracts correction signals $G = \{g_1, \dots, g_l\}$ encoding reusable constraints, such as masking fields for external recipients or confirming irreversible actions. Candidate updates ΔC_t may revise skill conditions, insert steps, prohibit actions, or require confirmation. A fixed scorer guides hill-climbing over textual skills, and the selected candidate enters a serial audit:

$$C_{t+1} = \text{AuditMerge}(C_t, \Delta C_t). \quad (9)$$

The audit has three gates. First, a schema and coverage gate rejects malformed skills, missing core conditions, and keyword-only rewrites that lack binding to an app, action, recipient, or sensitive evidence. Second, a fixed text scorer evaluates

rule coverage, held-out decisions, unsafe allows, and concise skill quality. Third, an LLM-Judge / sandbox gate audits confirmation timing, safety rationale, and executable state flow for downstream agents. This final gate catches skills that mention the right privacy rule and still produce unsafe action order or incomplete flow.

In our implementation, evolution is controlled hill climbing over text memory: mutations are accepted only when they improve the current best fixed-scorer skill and pass the serial audit. The process follows LLM-based textual optimization and recent skill-evolution systems (Yang et al., 2024; Jiang et al., 2026).

5 Experimental Setup

5.1 Research Questions and Evaluation Map

We organize evaluation around three experiment-level questions. E1 measures whether the local privacy pipeline can perceive sensitive evidence, improve retrieval-grounded arbitration, and construct SafeScreenshots. E2 diagnoses whether simpler deployment routes can replace local policy memory, and E3 evaluates whether user feedback can improve scenario-level skills. L1–L3 denote pipeline-

ID	Focus and evidence
E1	Local pipeline: L1 OCR coverage, L2 Open vs. Closed arbitration on 832 rows, and L3 SafeScreenshot audit.
E2	Policy memory: static, cloud, and edge-cloud routes under one evaluator.
E3	Skill evolution: five controlled scenarios with fixed scorer and sandbox gate.

Table 3: Experiment map. E1 combines component diagnostics with the benchmark L2 policy protocol. E2 uses the benchmark policy protocol. E3 evaluates controlled skill evolution.

layer diagnostics: L2 evaluates full-benchmark policy arbitration, while L1 and L3 audit visual triage, extraction, and SafeScreenshot construction. Server-side replay and L1/L3 details are reported in the appendices.

5.2 Compared Systems

We group baselines into detector, static-memory, cloud-reasoning, and routing classes (Siyan et al., 2025). Table 4 summarizes input context, raw-screen exposure, local policy access, and skill updates. E2 is a benchmark-wide diagnostic over 832 canonical rows, with systems varying exposure path, context, and access to policy memory. Cloud routes may see raw screenshots; expected labels, PII/risk tags, and policy-memory answers are withheld. Policy-grounded MaskClaw isolates static local policy memory, and skill-augmented MaskClaw is evaluated in E3.

5.3 Static Decision Protocol and Metrics

E1 and E2 use a static single-step protocol. Inputs include the screenshot, instruction, candidate intent, app/recipient context, and persona; expected labels, risk tags, and PII categories are evaluator-only. Systems output policy and route fields, and the evaluator recomputes the final Allow/Mask/Ask decision.

The decision governs exposure of the current visual observation. Downstream task success and action authorization are outside the scoring target. The evaluator applies a conservative priority: Ask overrides Mask, and Mask overrides Allow. Ask covers consent-sensitive or irreversible exposure, Mask covers redactable sensitive content, and Allow requires task necessity and persona-policy compatibility.

We report joint decision accuracy, Mask F1, Ask recall, leak rate, over-protection, and raw-upload rate. For error analysis, we define *Leak ...*, *Over-*

Protect ..., and *AskMiss ...*. Region-level masking errors, such as masking the wrong protected region, are treated as SafeScreenshot diagnostics and reported in Appendix D.

5.4 Layer-wise Diagnostics

E1 also includes layer-wise diagnostics. L2 uses all 832 policy rows; L1/L3 audit visual PII triage, sensitive-item extraction, and strict OCR-based SafeScreenshot construction on applicable subsets. Appendix D gives denominators, QA flags, and artifacts.

5.5 Self-Evolution Evaluation

E3 evaluates adaptation across five scenario families: iCloud cleanup, app permission, high-value transfer, schedule merge, and driving mode. Each scenario contains a task goal, three initial skills, user corrections, target constraints, and held-out decision and behavior tests. Corrections are evolution inputs; constraints and tests are fixed for scoring.

We use five scenarios, three start conditions, three random seeds, and 20 evolution iterations. Each skill is scored as:

$$J(\theta) = \lambda_1 \text{RuleCoverage} + \lambda_2 \text{TestDecisionAcc} + \lambda_3 \text{SafetyPenalty} + \lambda_4 \text{SkillQuality}, \quad (10)$$

where $\lambda_1 = 0.40$, $\lambda_2 = 0.30$, $\lambda_3 = 0.20$, and $\lambda_4 = 0.10$. Weights are fixed before evaluation. Rule coverage and held-out correctness dominate, safety is enforced by filters, and SkillQuality favors concise, inspectable skills. Rule-based caps or rejection handle malformed skills, critical safety violations, missing core rules, and unsafe allow decisions. The LLM-Judge audits privacy conditions, confirmation timing, and executable state transitions for downstream agents, following recent GUI evaluation practice (Zheng et al., 2023; Lin et al., 2025). We use BestScore@20 as the primary evolution metric and report behavior accuracy, unsafe action rate, and correction compliance as skill-use checks.

6 Results and Analysis

6.1 Local Arbitration and SafeScreenshot Checks

Table 5 shows the E1 evidence chain. The core L2 result improves policy accuracy from 0.533 ± 0.010 to 0.717 ± 0.007 . L1 reaches 0.967 normalized exact OCR text coverage, and L3 passes the strict no-

Category	System	Raw screen	Context	Local policy	Skill update
Detector	Static Regex / Presidio	No	none	No	No
Static memory	Static Policy Corpus	No	task/persona	static	No
Cloud VLM	Cloud VLM Minimal	Yes	minimal	No	No
	Cloud VLM Persona	Yes	persona	No	No
	Cloud Context-Complete Persona	Yes	full non-label	No	No
Routing	EdgeClaw-ClawXRouter	mixed	route features	No	No
Local ablation	MaskClaw Local VLM	No	full non-label	No	No
Ours	Policy-grounded MaskClaw	No	full non-label	static	No
	MaskClaw with Skills	No	full non-label	rules+skills	controlled

Table 4: Compared systems grouped by baseline class. Full non-label context excludes expected decisions, PII/risk tags, and policy-memory answers.

Check	Score	Mask R.	Allow R.	Ask R.
L1 OCR coverage	967	—	—	—
L2 Open	533±10	770±7	308±12	113±22
L2 Closed	717±7	804±6	683±18	377±32
L3 strict redaction	733	—	—	—

Table 5: E1 layer diagnostics. Values are reported in 10^{-3} : Score is normalized exact OCR text coverage for L1, policy accuracy for L2, and strict no-flag rate for L3. Only L2 Open and L2 Closed are directly comparable policy-decision results; component details are in Appendix D.

flag audit on most visible-PII cases. The L2 gain is contextual: Allow and Ask recall rise sharply while Mask recall remains high. Earlier frozen-interface diagnostics are reported in Appendix B.

6.2 Baseline Comparison

Static Regex is the strongest non-MaskClaw baseline by accuracy, reaching 0.557 despite having no task, persona, recipient, or policy memory. Visible-pattern detection provides a competitive lower bound and is close to the ungrounded local VLM ablation in Table 5. Its limited contextual policy semantics produce zero Ask recall, conflate task-necessary exposure with over-protection, and leave a leak rate of 0.283. Policy-grounded MaskClaw improves over this strongest baseline by 0.160 absolute accuracy, from 0.557 to 0.717. It also increases Mask F1 by 0.174 and reduces leak by 0.087.

Cloud and routing controls show different route-level behaviors. Cloud routes upload every raw screenshot, and Cloud Full-Context is dominated by Ask predictions, with high Ask recall but only 0.166 ± 0.005 accuracy. EdgeClaw-ClawXRouter reduces raw-cloud routing to 15.9%. As an




Figure 3: Skill-use checks in five controlled scenarios. The plot reports downstream behavior after text-level skill evolution.

exposure-routing proxy, it reaches 0.401 ± 0.006 accuracy. Table 6 shows that pattern detectors catch many visible spans, cloud routes tend to over-confirm or expose raw screens, and routing proxies leave final contextual arbitration unresolved. Appendix C reports rerun checks.

6.3 Skill Evolution and Behavior Improvement

Experiment 3 focuses on controlled skill evolution. The first question is whether correction-driven search can produce high-quality textual skills under a fixed scorer. The second is whether a separate sandbox gate rejects candidates whose text looks good and whose execution remains unsafe. Figure 3 checks whether selected skills also improve downstream behavior.

Across five controlled scenarios, BestScore@20 averages 88.76, with scores ranging from 77.70 to 93.67. The primary evolution result is that user corrections can be distilled into textual skills that satisfy fixed coverage, held-out decision, safety-penalty, and quality constraints in designed settings.

Method	Family	Acc.	Mask F1	Leak	Ask R.	Raw up.
Policy-grounded MaskClaw	policy memory	717 $\pm$ 7	819 $\pm$ 14	196 $\pm$ 6	377 $\pm$ 32	0
Static Regex	pattern detector	557 $\pm$ 0	645 $\pm$ 0	283 $\pm$ 0	0 $\pm$ 0	0
Static Policy Corpus	static memory	254 $\pm$ 0	413 $\pm$ 0	637 $\pm$ 0	438 $\pm$ 0	0
Cloud Minimal	cloud reasoning	308 $\pm$ 6	359 $\pm$ 14	709 $\pm$ 13	613 $\pm$ 45	1000
Cloud Persona	cloud reasoning	370 $\pm$ 4	472 $\pm$ 7	540 $\pm$ 9	263 $\pm$ 13	1000
Cloud Full-Context	cloud reasoning	166 $\pm$ 5	83 $\pm$ 6	956 $\pm$ 4	971 $\pm$ 19	1000
EdgeClaw-ClawXRouter [†]	exposure routing	401 $\pm$ 6	480 $\pm$ 3	552 $\pm$ 1	529 $\pm$ 19	159

Table 6: Architectural diagnostic of privacy routes. All values are reported in 10^{-3} . Values with $\pm$ summarize repeated runs; routing exposure is deterministic or route-derived. Policy-grounded MaskClaw includes local policy memory but excludes skill evolution; cloud routes are rule-free reasoning controls. Cloud Full-Context denotes Cloud Context-Complete Persona. Leak is expected Mask predicted as non-Mask. [†]EdgeClaw-ClawXRouter is evaluated as a route-level comparison; details are in Appendix C.




Figure 4: Post-optimization filtering from strict text scoring to the LLM-Judge / sandbox gate. Rejections mainly involve state flow, safety reasoning, and rule alignment.

Figure 3 shows that evolved skills improve held-out behavior decisions and reduce unsafe actions. Compliance remains imperfect (76.19% for transfer and 60.87% for driving), so a high-scoring textual skill can still miss evidence, wording, or order constraints when used by a downstream agent. This gap motivates the post-optimization sandbox filter below.

6.4 Sandbox Validation

Sandbox validation serves as the post-evolution safety audit. It checks confirmation timing, rationale, and executable state flow before a candidate can be treated as reusable memory.

Figure 4 reports post-optimization filtering from strict text scoring to the flow-level sandbox gate. The drop is expected because the two gates test different properties: iCloud moves from 66.14% strict pass to 13.23% sandbox pass, and transfer moves from 52.38% to 12.17%. Candidate skills can satisfy textual privacy constraints and still fail executable flow-level safety.

Thus, low sandbox pass rates show why deployment-oriented filtering is necessary after text-level skill search. The LLM-Judge provides an audit signal for reusable memory selection; full

counts and diagnostics are in the appendix.

6.5 Error Analysis

The remaining failures are primarily contextual boundary errors. In the 832-example joint-field diagnostic, the finalized arbitration interface reaches 0.7127 joint accuracy, with 593 correct cases and 239 joint errors. The three largest categories are protected content allowed, unnecessary masking, and confirmation replaced by masking, accounting for 79.5% of joint errors.

These errors cluster around persona, application, action, data type, and recipient boundaries: whether exposure is task-necessary, which recipient is trusted, and when confirmation is required. Retrieval diagnostics often find a related rule for the wrong boundary, suggesting boundary confusion and rule-applicability failures.

7 Conclusion

This paper frames GUI-agent screenshot privacy as exposure arbitration and presents MaskClaw, an edge-side framework that decides what should be allowed, masked, or confirmed under task, recipient, and persona context. This keeps privacy handling explicit while preserving useful visual context for downstream agents through allowed, masked, or confirmed observations at runtime. Across P-GUI-Evo, policy-memory-free or route-only baselines struggle with these boundaries and often expose raw screenshots. The skill-evolution study shows a narrower but useful result: in designed scenarios, user corrections can become reusable privacy skills, and sandbox checks catch candidates that still fail flow-level safety.

Limitations

P-GUI-Evo is a sanitized benchmark for GUI privacy decisions, built from reconstructed scenarios. It covers multiple personas, tasks, recipients, and UI perturbations, with personas serving as author-designed abstractions for workflow and policy variation. Real-user and organization-level validation remain future work. Our skill-evolution study is scenario-level: it shows that corrections can be distilled into reusable privacy skills in designed settings, while longer-term personalization requires broader user studies. The deployment scope is trusted edge-side execution in a user or organization environment such as a workstation or private server. These limits bound the empirical claims while preserving the core contribution: MaskClaw reduces raw screenshot exposure before downstream agents receive GUI observations.

Ethics Statement

P-GUI-Evo is built from simulated and sanitized GUI scenarios, with real user logs, private chats, medical records, payment credentials, identity documents, and personal identifiers outside the release scope. Released artifacts should be reviewed for accounts, contacts, organization identifiers, QR codes, credentials, and reverse-identifiable information; release-eligible assets should pass this review, while other assets should be withheld or replaced by metadata-only records. MaskClaw reduces raw screenshot exposure through local evidence extraction, policy arbitration, and visual redaction; cloud VLM and edge-cloud routes are diagnostic baselines. The intended use is GUI-agent privacy research, with high-risk medical, financial, identity, enterprise-compliance, and irreversible actions requiring separate professional or legal review. Evolved skills may encode user preferences and should remain inspectable; LLM judges, masking, and sandbox validation provide risk-reduction checks.

References

- Akari Asai, Zeqiu Wu, Yizhong Wang, Avi Sil, and Hannaneh Hajishirzi. 2024. Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection. *International Conference on Learning Representations*, 2024:9112–9141.
- Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. 2021. Extracting Training Data from Large Language Models. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 2633–2650.
- Zhiyuan Chang, Mingyang Li, Yuekai Huang, Ziyu Jiang, Xiaojun Jia, Qian Xiong, Junjie Wang, Zhaoyang Li, and Qing Wang. 2026. Know Thy Enemy: Securing LLMs Against Prompt Injection via Diverse Data Synthesis and Instruction-Level Chain-of-Thought Learning. *Preprint*, arXiv:2601.04666.
- Karmesh Siddharam Chaudhari, Youxiang Zhu, Amy Feng, Xiaohui Liang, and Honggang Zhang. 2026. Assessing Privacy Preservation and Utility in Online Vision-Language Models. *Preprint*, arXiv:2604.09695.
- Yining Chen, Jihao Zhao, Bo Tang, Haofen Wang, Yue Zhang, Fei Huang, Feiyu Xiong, and Zhiyu Li. 2026. MemPrivacy: Privacy-Preserving Personalized Memory Management for Edge-Cloud Agents. *Preprint*, arXiv:2605.09530.
- Qingyan Guo, Rui Wang, Junliang Guo, Bei Li, Kaitao Song, Xu Tan, Guoqing Liu, Jiang Bian, and Yujie Yang. 2024. Connecting Large Language Models with Evolutionary Algorithms Yields Powerful Prompt Optimizers. *International Conference on Learning Representations*, 2024:34133–34156.
- Ivan Habernal, Fatemehsadat Mirehghallah, Patricia Thaine, Sepideh Ghanavati, and Oluwaseyi Feyisetan. 2023. Privacy-Preserving Natural Language Processing. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics: Tutorial Abstracts*, pages 27–30, Dubrovnik, Croatia. Association for Computational Linguistics.
- Wenyi Hong, Weihang Wang, Qingsong Lv, Jiazheng Xu, Wenmeng Yu, Junhui Ji, Yan Wang, Zihan Wang, Yuxiao Dong, Ming Ding, and Jie Tang. 2024. CogAgent: A Visual Language Model for GUI Agents. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14281–14290.
- Shengding Hu, Yuge Tu, Xu Han, Chaoqun He, Ganqu Cui, Xiang Long, Zhi Zheng, Yewei Fang, Yuxiang Huang, Weilin Zhao, Xinrong Zhang, Zheng Leng Thai, Kaihuo Zhang, Chongyi Wang, Yuan Yao, Chenyang Zhao, Jie Zhou, Jie Cai, Zhongwu Zhai, and 6 others. 2024. MiniCPM: Unveiling the Potential of Small Language Models with Scalable Training Strategies. *arXiv preprint*.
- Deyang Jiang, Jing Huang, Xuanle Zhao, Lei Chen, Liming Zheng, Fanfan Liu, Haibo Qiu, Peng Shi, and Zhixiong Zeng. 2026. TreeCUA: Efficiently Scaling GUI Automation with Tree-Structured Verifiable Evolution. <https://arxiv.org/abs/2602.09662v1>.
- Siwon Kim, Sangdoo Yun, Hwaran Lee, Martin Gubri, Sungroh Yoon, and Seong Joon Oh. 2023. ProPILE:

- Probing Privacy Leakage in Large Language Models. *Advances in Neural Information Processing Systems*, 36:20750–20762.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474. Curran Associates, Inc.
- Kevin Qinghong Lin, Siyuan Hu, Linjie Li, Zhengyuan Yang, Lijuan Wang, Philip Torr, and Mike Zheng Shou. 2025. [Computer-Use Agents as Judges for Generative User Interface](#). *arXiv preprint*.
- Xiao Liu, Bo Qin, Dongzhu Liang, Guang Dong, Hanyu Lai, Hanchen Zhang, Hanlin Zhao, Iat Long Iong, Jiadai Sun, Jiaqi Wang, Junjie Gao, Junjun Shan, Kangning Liu, Shudan Zhang, Shuntian Yao, Siyi Cheng, Wentao Yao, Wenyi Zhao, Xinghan Liu, and 11 others. 2024. Auto-GLM: Autonomous Foundation Agents for GUIs. <https://arxiv.org/abs/2411.00820v1>.
- Yadong Lu, Jianwei Yang, Yelong Shen, and Ahmed Awadallah. 2024a. [OmniParser for Pure Vision Based GUI Agent](#). *Preprint*, arXiv:2408.00203.
- Yujie Lu, Dongfu Jiang, Wenhui Chen, William Yang Wang, Yejin Choi, and Bill Yuchen Lin. 2024b. [Wild-Vision: Evaluating Vision-Language Models in the Wild with Human Preferences](#). *Advances in Neural Information Processing Systems*, 37:48224–48255.
- Ziyu Ma, Shidong Yang, Yuxiang Ji, Xucong Wang, Yong Wang, Yiming Hu, Tongwen Huang, and Xiangxiang Chu. 2026. [SkillClaw: Let Skills Evolve Collectively with Agentic Evolver](#). *Preprint*, arXiv:2604.08377.
- Helen Nissenbaum. 2004. Privacy as Contextual Integrity. *Washington Law Review*, 79(1):119.
- Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. 2023. [MemGPT: Towards LLMs as Operating Systems](#). *arXiv preprint*.
- Zehan Qi, Xiao Liu, Iat Long Iong, Hanyu Lai, Xueqiao Sun, Jiadai Sun, Xinyue Yang, Yu Yang, Shuntian Yao, Wei Xu, Jie Tang, and Yuxiao Dong. 2025. WebRL: Training LLM Web Agents via Self-Evolving Online Curriculum Reinforcement Learning. *International Conference on Learning Representations*, 2025:79791–79821.
- Christopher Rawles, Alice Li, Daniel Rodriguez, Oriana Riva, and Timothy Lillicrap. 2023. AndroidInTheWild: A Large-Scale Dataset For Android Device Control. *Advances in Neural Information Processing Systems*, 36:59708–59728.
- Alireza Salemi, Sheshera Mysore, Michael Bendersky, and Hamed Zamani. 2023. [LaMP: When Large Language Models Meet Personalization](#). *arXiv preprint*.
- Noah Shinn, Federico Cassano, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. 2023. Reflexion: Language agents with verbal reinforcement learning. *Advances in Neural Information Processing Systems*, 36:8634–8652.
- Li Siyan, Vethavikashini Chithra Raghuram, Omar Khattab, Julia Hirschberg, and Zhou Yu. 2025. [PA-PILLON: Privacy Preservation from Internet-based and Local Language Model Ensembles](#). In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 3371–3390, Albuquerque, New Mexico. Association for Computational Linguistics.
- Batuhan Tömekçe, Mark Vero, Robin Staab, and Martin Vechev. 2024. [Private Attribute Inference from Images with Vision-Language Models](#). *Advances in Neural Information Processing Systems*, 37:103619–103651.
- Junyang Wang, Haiyang Xu, Jiabo Ye, Ming Yan, Weizhou Shen, Ji Zhang, Fei Huang, and Jitao Sang. 2024. [Mobile-Agent: Autonomous Multi-Modal Mobile Device Agent with Visual Perception](#). *arXiv preprint*.
- Shouju Wang, Fenglin Yu, Xirui Liu, Xiaoting Qin, Jue Zhang, Qingwei Lin, Dongmei Zhang, and Saravan Rajmohan. 2025a. [Privacy in Action: Towards Realistic Privacy Mitigation and Evaluation for LLM-Powered Agents](#). *Preprint*, arXiv:2509.17488.
- Shuai Wang, Weiwen Liu, Jingxuan Chen, Yuqi Zhou, Weinan Gan, Xingshan Zeng, Yuhua Che, Shuai Yu, Xinlong Hao, Kun Shao, Bin Wang, Chuhan Wu, Yasheng Wang, Ruiming Tang, and Jianye Hao. 2025b. [GUI Agents with Foundation Models: A Comprehensive Survey](#). *Preprint*, arXiv:2411.04890.
- Yanxi Wang, Zhiling Zhang, Wenbo Zhou, Weiming Zhang, Jie Zhang, Qiannan Zhu, Yu Shi, Shuxin Zheng, and Jiyan He. 2026a. [GUIGuard-Bench: Toward a General Evaluation for Privacy-Preserving GUI Agents](#). *arXiv preprint*.
- Yidan Wang, Yanan Cao, Yubing Ren, Fang Fang, Zheng Lin, and Binxing Fang. 2025c. [PIG: Privacy Jailbreak Attack on LLMs via Gradient-based Iterative In-Context Optimization](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 9645–9660, Vienna, Austria. Association for Computational Linguistics.
- Youting Wang, Yuan Tang, Yitian Qian, and Chen Zhao. 2026b. [VisualLeakBench: Auditing the Fragility of Large Vision-Language Models against PII Leakage and Social Engineering](#). *Preprint*, arXiv:2603.13385.

- Peng Xia, Jianwen Chen, Hanyang Wang, Jiaqi Liu, Kaide Zeng, Yu Wang, Siwei Han, Yiyang Zhou, Xujing Zhao, Haifeng Chen, Zeyu Zheng, Cihang Xie, and Huaxiu Yao. 2026. [SkillRL: Evolving Agents via Recursive Skill-Augmented Reinforcement Learning](#). *arXiv preprint*.
- Yifan Xu, Xiao Liu, Xueqiao Sun, Siyi Cheng, Hao Yu, Hanyu Lai, Shudan Zhang, Dan Zhang, Jie Tang, and Yuxiao Dong. 2025. [AndroidLab: Training and Systematic Benchmarking of Android Autonomous Agents](#). In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2144–2166, Vienna, Austria. Association for Computational Linguistics.
- Haolong Yan, Jia Wang, Xin Huang, Yeqing Shen, Ziyang Meng, Zhimin Fan, Kaijun Tan, Jin Gao, Lieyu Shi, Mi Yang, Shiliang Yang, Zhirui Wang, Brian Li, Kang An, Chenyang Li, Lei Lei, Mengmeng Duan, Danxun Liang, Guodong Liu, and 80 others. 2025. [Step-GUI Technical Report](#). *Preprint*, arXiv:2512.15431.
- Chengrun Yang, Xuezhi Wang, Yifeng Lu, Hanxiao Liu, Quoc V. Le, Denny Zhou, and Xinyun Chen. 2024. Large Language Models as Optimizers. *International Conference on Learning Representations*, 2024:12028–12068.
- Jianwei Yang, Hao Zhang, Feng Li, Xueyan Zou, Chunyuan Li, and Jianfeng Gao. 2023. Set-of-Mark Prompting Unleashes Extraordinary Visual Grounding in GPT-4V. <https://arxiv.org/abs/2310.11441v2>.
- Yuan Yao, Tianyu Yu, Ao Zhang, Chongyi Wang, Junbo Cui, Hongji Zhu, Tianchi Cai, Haoyu Li, Weilin Zhao, Zhihui He, Qianyu Chen, Huarong Zhou, Zhensheng Zou, Haoye Zhang, Shengding Hu, Zhi Zheng, Jie Zhou, Jie Cai, Xu Han, and 4 others. 2024. [MiniCPM-V: A GPT-4V Level MLLM on Your Phone](#). *arXiv preprint*.
- Keen You, Haotian Zhang, Eldon Schoop, Floris Weers, Amanda Swearngin, Jeffrey Nichols, Yinfei Yang, and Zhe Gan. 2024. [Ferret-UI: Grounded Mobile UI Understanding with Multimodal LLMs](#). *Preprint*, arXiv:2404.05719.
- Mert Yuksekgonul, Federico Bianchi, Joseph Boen, Sheng Liu, Zhi Huang, Carlos Guestrin, and James Zou. 2024. [TextGrad: Automatic "Differentiation" via Text](#). *Preprint*, arXiv:2406.07496.
- Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah Goodman. 2022. STaR: Bootstrapping Reasoning With Reasoning. *Advances in Neural Information Processing Systems*, 35:15476–15488.
- Chi Zhang, Zhao Yang, Jiaxuan Liu, Yanda Li, Yucheng Han, Xin Chen, Zebiao Huang, Bin Fu, and Gang Yu. 2025. [AppAgent: Multimodal Agents as Smartphone Users](#). In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems*, CHI '25, pages 1–20, New York, NY, USA. Association for Computing Machinery.
- Jie Zhang, Xiangkui Cao, Zhouyu Han, Shiguang Shan, and Xilin Chen. 2026. [Multi-PA: A Multi-perspective Benchmark on Privacy Assessment for Large Vision-Language Models](#). *Preprint*, arXiv:2412.19496.
- Lepeng Zhao, Zhenhua Zou, Shuo Li, and Zhuotao Liu. 2026. Anonymization-Enhanced Privacy Protection for Mobile GUI Agents: Available but Invisible. <https://arxiv.org/abs/2602.10139v3>.
- Boyuan Zheng, Boyu Gou, Jihyung Kil, Huan Sun, and Yu Su. 2024. [GPT-4V\(ision\) is a Generalist Web Agent, if Grounded](#). *Preprint*, arXiv:2401.01614.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, Hao Zhang, Joseph Gonzalez, and Ion Stoica. 2023. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. *Advances in Neural Information Processing Systems*, 36:46595–46623.
- Huichi Zhou, Siyuan Guo, Anjie Liu, Zhongwei Yu, Ziqin Gong, Bowen Zhao, Zhixun Chen, Menglong Zhang, Yihang Chen, Jinsong Li, Runyu Yang, Qiangbin Liu, Xinlei Yu, Jianmin Zhou, Na Wang, Chunyang Sun, and Jun Wang. 2026. Memento-Skills: Let Agents Design Agents. <https://arxiv.org/abs/2603.18743v1>.
- Youxiang Zhu, Ning Gao, Xiaohui Liang, and Honggang Zhang. 2024. [Exploiting Privacy Preserving Prompt Techniques for Online Large Language Model Usage](#). In *GLOBECOM 2024 - 2024 IEEE Global Communications Conference*, volume 2024, pages 4304–4309.

Partition	Count	Share
<i>Policy label</i>		
Mask	438	52.6%
Allow	314	37.7%
Ask	80	9.6%
<i>Persona</i>		
User A	234	28.1%
User B	246	29.6%
User C	352	42.3%
<i>Bucket</i>		
D1 basic	174	20.9%
D2 generalization	546	65.6%
D3 stress	112	13.5%
<i>Scenario variant</i>		
Base screenshot	250	30.0%
Structural/UI shift (ss)	201	24.2%
Visual shift (vs)	208	25.0%
Lexical/semantic shift (1s, lower-case L)	173	20.8%

Table 7: Dataset distribution used by the static policy-decision experiments. Percentages are computed over 832 samples. Scenario variants preserve the same underlying privacy boundary while changing the UI structure, visual rendering, or task wording.

A Evaluation Protocol and Dataset Details

Prediction methods receive only fields that would be available to a deployed privacy mediator: the screenshot, application context, persona context, user instruction, and candidate agent intent. Audit-only fields, including PII category, risk tags, expected labels, and author audit notes, are used only for offline scoring and diagnostics. This separation prevents the evaluator from leaking expected policy information into the decision process.

P-GUI-Evo contains 832 controlled GUI privacy samples built from reconstructed UI layouts, LLM-assisted scenario drafting, typed sanitized sensitive-field injection, and audit filtering. Tables and prose use the display labels Allow (314), Mask (438), and Ask (80). The benchmark includes three personas and three buckets, summarized in Table 7. D1 contains clean in-distribution cases, D2 contains UI and task variants, and D3 contains stress cases with boundary conditions and UI perturbations. The final set also keeps controlled variants around shared scenario cores: structural UI shifts (ss), visual shifts (vs), and lexical or semantic task shifts (1s, lowercase L-s).

B L2 Retrieval-Augmented Arbitration Diagnostics

The L2 Closed result in the main text reports the finalized retrieval-augmented policy arbitration in-

terface on the canonical 832-sample benchmark. We include an earlier interface only to show how the final design changes the error profile; the main claims use the finalized interface.

Table 8 shows that relevant rules are usually present: the full-set top-k oracle reaches 0.9940. The remaining gap is mainly an arbitration gap, where related rules are retrieved but applied to the wrong policy boundary.

Split	N	Policy	Mask	Ask-R
Preliminary pilot	120	758	900	700
Preliminary held-out	732	616	824	478
Preliminary full set	832	609	821	400

Table 8: Preliminary-interface diagnostics across pilot, held-out, and full splits. Metric values are reported in 10^{-3} . These diagnostics are included only to contextualize the finalized interface.

Interface	Policy	Mask R.	Mask Acc.	Ask R.
Preliminary interface	609	821	722	400
Finalized interface	717 $\pm$ 7	784 $\pm$ 15	904 $\pm$ 3	377 $\pm$ 32

Table 9: Preliminary versus finalized retrieval-augmented policy arbitration on the full 832-sample set. Values are reported in 10^{-3} . The finalized interface is the setting used for the main result.

Mediator responsibility. The L2 Closed setting uses a deterministic mediator to choose among top-k rules. The closed pipeline is:

query/context $\rightarrow$ local rule retrieval $\rightarrow$ deterministic mediator $\rightarrow$ local VLM visual evidence check.

The deterministic mediator sets the selected rule id, mask flag, and confirmation flag. The local VLM checks whether the screenshot contains supporting visual evidence or an obvious contradiction. The mediator owns the locked policy fields, while the VLM reports visual evidence and contradiction flags. This responsibility split attributes the finalized-interface gain to rule-grounded arbitration with retrieval support and visual verification.

C Full Baseline Metrics and Route-Level Comparison

The main text reports a route-oriented architectural diagnostic; Table 10 gives the full score table. All systems are rescored against the canonical 832-row expected-decision table. Policy-grounded MaskClaw isolates rule-grounded static policy mediation, while cloud routes are rule-free controls. Skill recovery and skill-routing experiments are excluded from this diagnostic.

Method	Acc.	M-P	M-R	M-F1	Leak	FPR	Ask-R	Route	Raw	Masked	Local
Policy-grounded MaskClaw	717	836	804	819	196	244	377	local	0	0	1000
Static Regex	557	587	717	645	283	525	0	local	0	0	1000
Static Policy Corpus	254	480	363	413	637	405	438	local	0	0	1000
Cloud Minimal	308	471	294	359	709	385	613	cloud	1000	0	0
Cloud Persona	370	490	470	472	540	535	263	cloud	1000	0	0
Cloud Context-Complete Persona	166	724	48	83	956	22	971	cloud	1000	0	0
EdgeClaw-ClawXRouter	401	512	448	480	552	484	529	hybrid	159	460	381

Table 10: Full route-diagnostic metrics on the 832-sample benchmark with explicit routing exposure. Values are reported in 10^{-3} . Main policy metrics follow the rerun summary where available; routing columns report exposure-mode rates on the same scale.

Repeated-run stability. Repeated runs on the same 832 canonical samples show the same pattern: policy-grounded MaskClaw remains the strongest local policy route, cloud routes remain conservative and Ask-heavy, and EdgeClaw-ClawXRouter is best interpreted as a route-level comparison rather than a contextual arbitration system.

Method	Acc.	Mask F1	Leak	Ask-R
Regex	557 $\pm$ 0	645 $\pm$ 0	283 $\pm$ 0	0 $\pm$ 0
Policy-grounded	717 $\pm$ 7	819 $\pm$ 14	196 $\pm$ 6	377 $\pm$ 32
Static Policy	254 $\pm$ 0	413 $\pm$ 0	637 $\pm$ 0	438 $\pm$ 0
Cloud Min.	308 $\pm$ 6	359 $\pm$ 14	709 $\pm$ 13	613 $\pm$ 45
Cloud Pers.	370 $\pm$ 4	472 $\pm$ 7	540 $\pm$ 9	263 $\pm$ 13
Cloud Full	166 $\pm$ 5	83 $\pm$ 6	956 $\pm$ 4	971 $\pm$ 19
EdgeClaw	401 $\pm$ 6	480 $\pm$ 3	552 $\pm$ 1	529 $\pm$ 19

Table 11: Three-run stability checks for valid repeated controls. Values are reported in 10^{-3} as mean $\pm$ sample standard deviation on the same 832 canonical samples.

Cloud Context-Complete Persona is a rule-free cloud control. It sees the raw screenshot plus app, command, agent intent, recipient context, and persona profile; expected labels, PII/risk tags, policy-memory answers, and label-revealing identifiers are withheld. Its high leak value mostly comes from expected Mask samples predicted as Ask, with few direct Allow predictions, indicating over-confirmation as the dominant behavior.

Expected	Pred Allow	Pred Mask	Pred Ask
Allow	44	7	263
Mask	2	21	415
Ask	3	1	76

Table 12: Cloud Context-Complete Persona confusion summary. The baseline predicts Ask for 754/832 samples, indicating over-confirmation as its dominant arbitration behavior.

Table 12 explains the low Cloud Full-Context score. The model rarely directly allows protected cases, but it routes most policy decisions to confirmation. This conservative behavior is safer than unconstrained Allow, yet it fails the benchmark’s

policy-arbitration objective because many cases require a task-specific Mask or Allow decision.

EdgeClaw-ClawXRouter emits routing states for local, masked-cloud, and raw-cloud exposure. Its safety objective differs from P-GUI-Evo: it decides whether an observation should remain local, be masked before cloud use, or be sent raw, whereas our benchmark asks for a contextual exposure policy. Table 14 maps its routes to the nearest Allow/Mask/Ask labels only for route-level comparison. In the original routing output, 46.0% of samples are masked before cloud upload and 38.1% remain local.

D L1/L3 Component Checks and SafeScreenshot Audit

The L1 and L3 checks are module-level diagnostics. L1 uses visual triage and sensitive-item extraction, and the main-table L1 score reports OCR text coverage by matching normalized gold sensitive-text values against OCR output. L3 evaluates automatic SafeScreenshot construction: sensitive items are extracted, OCR localizes matching text, and a strict remasking rule masks matched regions while producing conservative audit flags.

Table 16 reports the strict L3 run used in the main text. It matches OCR boxes to extracted sensitive-item text or token-like values; broad type-keyword expansion and global regex fallback are used only in the recall-first diagnostic variant. The no-flag count is an automatic audit status for clear leakage or overmasking checks. Manual inspection verified that the strict run removed the broad overmasking patterns observed in the recall-first run.

E Skill Evolution Scenario Protocol and Full Results

Experiment 3 evaluates scenario-level adaptation. A scenario includes a task goal, initial skills, user corrections, target constraints, and held-out behav-

Cloud baseline	Image	Temp.	Context	Output handling
Minimal	raw	0.0	app + command	parsed JSON policy fields
Persona	raw	0.0	app + command + persona profile	parsed JSON policy fields
Context-Complete Persona	raw	0.0	app; command; intent; recipient context; persona	parsed JSON policy fields

Table 13: Cloud baseline configuration. The table reports the information exposed to each rule-free cloud control; all expected labels and analysis-only metadata are withheld.

State	Route label	Upload mode	Mapped decision	Count	Share
S1	cloud_raw	raw_image	Allow	132	15.9%
S2	cloud_after_mask	masked_image	Mask	383	46.0%
S3	local_only	local_only	Ask	317	38.1%

Table 14: Route-level comparison for EdgeClaw-ClawXRouter. Counts sum to the 832-sample benchmark. The mapped decision is used only to place the router in the shared policy-decision table; the upload mode preserves the hybrid edge-cloud behavior.

Check	N	Result
Visual triage	832	832 ok; 620 visible-PII
L1 OCR text coverage	736	1,446/1,496 exact; 99.96% char recall
PII-type audit	251	reviewed for PII type
Visible-PII L3 set	558	466 mask + 92 allow
Sensitive-item extraction	558	558 ok; zero empty rows

Table 15: L1/L3 diagnostic denominators. These audits support benchmark checking as component-level quality signals.

Strict SafeScreenshot audit	Count
Rows	558
Mask generated	475
Automatic no-flag cases	409
Flagged for targeted audit	149
Missed sensitive-token flags	145
No OCR match flags	83
Overmask-suspect flags	5

Table 16: Strict SafeScreenshot construction audit on 558 visible-PII screenshots. Flags are conservative automatic audit signals and may overlap.

ior tests. The update target is a validated skill for downstream behavior; the static policy memory used for arbitration remains fixed in this experiment.

These results support controlled behavior improvement in five designed scenarios. The claim scope is scenario-level adaptation under fixed tests and curated feedback signals. In particular, compliance below 100% in the transfer and driving scenarios means a behavior can become correct at a high level while still missing required evidence, confirmation timing, or ordering constraints.

Variant	Mask generated	No-flag cases
Recall-first	542/558	338/558
Strict	475/558	409/558

Table 17: Recall-first versus strict remasking. Recall-first masks more text but triggers more audit flags; strict is the reported setting.

F Sandbox and LLM-Judge Validation Details

The strict scorer checks text-level skill quality and rule coverage. The LLM-Judge / sandbox gate adds privacy-condition, confirmation-timing, safety-reasoning, and executable-flow checks.

The gap between strict and sandbox pass rates means the sandbox gate rejects candidates that look acceptable as text but fail flow-level constraints. Common failure modes include missing state transitions, unsafe confirmation ordering, and weak alignment with the scenario’s target constraints.

G Error Analysis and Retrieval Diagnostics

The main residual failures are contextual boundary errors. In the 832-example joint-field diagnostic, L2 Closed reaches 0.7127 joint accuracy, with 593 correct cases and 239 joint errors. The most common error categories are protected content allowed, unnecessary masking, and confirmation replaced by masking, accounting for 79.5% of joint errors.

The equality between 239 wrong-boundary retrieval cases and 239 joint errors is numerical. A wrong-boundary retrieval means the retrieved rule is topically related but aligned to the wrong policy boundary. This supports the paper’s main interpretation: static rules require reliable applicability judgment and execution gates.

Scenario	Task boundary	Optimization	Tests	Evaluation note
iCloud cleanup	Avoid deleting unsynced, favorited, or irreversible local media.	9 starts, 20 iters	8	Full controlled setting.
App permission	Replace automatic permission approval with minimal, contextual authorization.	9 starts, 20 iters	8	Full controlled setting.
High-value transfer	Enforce amount thresholds, liveness, payee checks, and final confirmation.	9 starts	15	Valid-run subset with fixed behavior tests.
Calendar merge	Merge calendars while respecting lunch blocks, buffers, attendees, time zones, and confirmation.	9 starts, 20 iters	8	Full controlled setting.
Driving mode	Minimize driver attention and prevent sensitive or non-urgent message disclosure.	9 starts	8	Diagnostic subset with fixed behavior tests.

Table 18: Skill-evolution protocol by scenario. Evaluation notes distinguish full controlled settings from valid diagnostic subsets.

Scenario	Best@20	Tests	Base Acc.	Evo Acc.	Base Unsafe	Evo Unsafe	Compliance
iCloud photo cleanup	86.67	8	37.50	100.00	62.50	12.50	100.00
App permission	93.67	8	25.00	100.00	75.00	0.00	100.00
High-value transfer	92.96	15	20.00	100.00	80.00	0.00	76.19
Calendar merge	77.70	8	12.50	100.00	87.50	0.00	95.65
Driving mode	92.78	8	25.00	100.00	75.00	0.00	60.87
Macro average	88.76	–	24.00	100.00	76.00	2.50	–

Table 19: Full behavior results for the five controlled skill-evolution scenarios. Best@20 is the mean best strict score within 20 iterations where applicable. Accuracy, unsafe action rate, and correction compliance are distinct metrics.

Scenario	N	Strict	Judge	Sandbox	Reject
iCloud	189	66.14	14.29	13.23	80.00
App permission	189	46.03	39.15	39.15	14.94
Transfer	189	52.38	12.70	12.17	76.77
Calendar	189	43.39	20.11	16.93	60.98
Driving	9	100.00	33.33	33.33	66.67

Table 20: Sandbox validation pass rates. Values are percentages except candidate count. Driving uses a smaller diagnostic candidate set.

The skill-recovery results serve as a boundary and future-work signal. Reliable applicability judgment and an execution gate are required before targeted skills can be merged into a general policy system.

H Prompt, Mediator, and Judge Templates

We include canonical templates as a compact summary of the prompting and mediation interfaces. The templates show the responsibilities, fields, and output contracts needed for benchmark-level reproduction while keeping sensitive release details, author-identifying metadata, and sample-specific expansions out of the paper. Complete sanitized prompts and mediation code belong in the released artifact.

L2 Open policy template. The L2 Open baseline receives the raw screenshot and task context; retrieved rules and mediator fields are withheld. Its

role is to predict the two canonical decision fields. The instructions frame the task as contextual exposure arbitration: visible sensitive-looking content triggers masking only when policy context requires it, task need triggers confirmation only under the relevant consent boundary, and Ask is reserved for confirmation-sensitive cases. The evaluator derives the final policy label from the two canonical fields.

Role: on-device privacy policy checker for whether the original screenshot can be exposed to an untrusted downstream agent under local protection policy.

Inputs: app/platform, persona, relationship tag, user task, and screenshot.

Decision fields: mask flag and confirmation flag.

Confirmation trigger: use confirmation only when recipient identity, authorization boundary, identifier semantics, or visual evidence is unclear.

Required JSON: mask flag, confirmation flag, and rationale.

L2 Closed mediator-locked template. The L2 Closed setting inserts retrieval and a deterministic mediator before the local VLM visual check. The retrieved policy rules and mediator decision are rendered into the prompt, and the mediator owns the final policy fields. The VLM checks whether the screenshot contains visual evidence that supports the mediator decision or a clear visual contradiction. When it detects a contradiction, it reports the contradiction while preserving the locked policy fields in the same response.

Role: visual-evidence and contradiction checker under mediator-locked policy fields.

Scenario	Strict	Gate	Judge	Rule	Safety	Flow	Other
iCloud	64	62	84	93	80	161	4 other
App permission	102	102	104	103	114	115	1 keyword
Transfer	90	90	118	97	134	165	2 keyword
Calendar	107	50	106	116	133	146	1 other
Driving	0	0	4	4	6	6	0

Table 21: Major sandbox and judge rejection signals. Counts are diagnostic flags and can overlap because a candidate can trigger multiple failure reasons. The recurring state-flow and safety failures explain why sandbox validation is stricter than BestScore@20 or strict text scoring alone.

Joint-field error	Count	Share of errors
Protected content allowed	95	39.7%
Unnecessary masking	58	24.3%
Confirmation replaced by masking	37	15.5%
Other joint-field errors	49	20.5%
Total joint errors	239	100.0%

Table 22: Joint-field error categories for L2 Closed. The top three categories account for 79.5% of all joint errors and mostly reflect contextual boundary decisions.

Retrieval diagnostic	Count
Wrong-boundary retrieval	239
Weak retrieval evidence	8
Neither	585

Table 23: Retrieval diagnostics for L2 Closed. Wrong-boundary retrieval is a diagnostic category for policy-boundary confusion.

Inputs: app/platform, persona context, relationship context, user instruction, agent intent, top-k policy rules, mediator-selected rule id, mediator mask flag, mediator confirmation flag, mediator policy decision, and mediator decision basis.

Locked fields: selected rule id, mask flag, and confirmation flag must be copied from the mediator values.

Visual check fields: visual evidence, visual contradiction flag, and rationale.

Required JSON: selected rule id, mask flag, confirmation flag, rationale, visual evidence, and contradiction flag.

Field group	Source and overwrite boundary
Retrieved rules	topk_rules_block; retrieval context consumed by the deterministic mediator.
Locked policy fields	selected rule id, mask flag, confirmation flag, and derived policy decision; copied by the local VLM and used for scoring.
VLM diagnostics	visual evidence and visual contradiction flag; generated by the local VLM but unable to overwrite policy fields.

Table 25: Controlled L2 policy interface. The local VLM adds visual support or contradiction diagnostics while the mediator keeps the policy fields fixed for the benchmark run.

Deterministic mediator boundary. The reported L2 Closed result uses a deterministic mediator instance over structured retrieval-query facets and top-k rule metadata. The mediator maps the retrieved rule context and scenario facets to a selected rule id, a mask flag, a confirmation flag, and

a derived policy decision. The implementation uses a deterministic rule-priority and fallback procedure. This design separates variable policy knowledge from a fixed execution interface.

The implementation is treated as experiment-reproduction code. In the paper, we report its interface and reproduction boundary because the claim is about benchmark-level retrieval-augmented policy arbitration. The final L2 Closed setting reaches 0.717 ± 0.007 policy accuracy on the 832-sample benchmark.

Rule and skill terminology. Rules support policy arbitration in the static L2 mediator. Skills are validated behavior policies used in the self-evolution setting. Feedback may produce candidate constraints, rule-like conditions, or skill patches, and the Experiment 3 claim focuses on behavior-improving skills in controlled scenarios.

Judge and sandbox rubric. The judge checks whether a candidate skill specifies the relevant privacy condition, asks for confirmation at the right time, avoids unsafe action defaults, and provides executable state-flow logic. A pass means the candidate satisfies the audit protocol for that controlled scenario.

I Implementation Boundary

The reported experiments evaluate offline policy arbitration, component-level SafeScreenshot construction, and controlled text-skill adaptation. We do not claim end-to-end device latency or deployment readiness. Live deployment would additionally require screenshot capture, model loading, network variability for cloud controls, GUI execution, and user-confirmation latency.

J Annotation and Scenario Construction

The benchmark scenarios are reconstructed and sanitized. Scenario construction records the persona abstraction, app/platform, user instruction, candidate agent action, recipient or relationship context, sensitive evidence, and expected policy decision.

Exploratory setting	Recovery	Pres. harm	OOS harm	Net	Use
Routed skill application	0	–	–	–	diagnostic only
Sampled targeted skills	5	2	5	-2	limitation

Table 24: Targeted skill-recovery diagnostics. Targeted skills can recover some cases, but preservation and out-of-scope harms prevent using them as a main full-set performance result.

Labels distinguish whether the raw screenshot can be exposed, whether a SafeScreenshot is required, or whether the system should pause for confirmation. The construction uses author-audited benchmark labels for controlled evaluation, with real user logs, real private chats, external crowdworkers, and human-subject interactions outside the data source.

For Experiment 3, each controlled scenario defines task goals, initial skills, user corrections, distilled constraints, and held-out behavior tests. Candidate skills are evaluated under fixed scoring and sandbox checks. This construction supports controlled adaptation claims over designed feedback signals.

Step	Recorded fields and validation boundary
UI pattern abstraction	App/platform family, interaction template, layout class, and task surface; patterns approximate common GUI interactions through reconstructed layouts.
LLM-assisted scenario drafting	Persona abstraction, task instruction, candidate action, and recipient or relationship context; drafts are prompt generated and sanitized.
Policy-label consistency pass	Allow, Mask, Ask, lower-level mask and confirmation fields; inconsistent records are rejected or repaired before scoring.
Sensitive-field injection	Typed fake values, visual/textual evidence, localizable-region status, PII/risk category, and benchmark bucket; audit fields stay offline.
HTML/UI rendering	Reconstructed layout, render viewport, and screenshot path; rendered screenshots become benchmark and SafeScreenshot inputs.
Label correction and audit	Canonical expected-decision table, correction notes, and E3 repair records; known corrections are synchronized before scoring.
Skill scenario construction	Task goal, initial skills, correction signals, distilled constraints, and held-out tests; used only for controlled adaptation experiments.

Table 26: Scenario and label construction protocol. The table records what is controlled during benchmark construction and what remains offline-only audit information.

Sanitized instruction excerpt. Scenario construction instructions require the project authors to specify: (i) the persona and privacy sensitivity, (ii) the application and recipient boundary, (iii) the candidate agent action, (iv) visible sensitive evidence and whether it is localizable, (v) the expected policy decision, and (vi) a short rationale explaining why the current GUI state should be exposed, masked, or paused for confirmation. For skill-evolution scenarios, the corresponding record also includes user corrections, target constraints, and held-out behavior tests. Because no real users or external annotators are involved, there is no par-

ticipant compensation or consent process for the dataset construction reported here.

K Artifact, Model, and Reproducibility Details

Experiment	Inputs, controls, and boundary
E1 L2	832 rows, prompts, policy memory, predictions, and fixed evaluator.
E1 L1/L3	832 triage rows, 736 OCR-coverage rows, and 558 visible-PII SafeScreenshot rows; component diagnostics only.
E2 baseline	832 rows and prediction files, recomputed against expected labels; routing systems use the route-level mapping in Appendix C.
E3 skills	Five scenario specs, correction signals, candidate skills, behavior tests, and recorded seeds.
Sandbox	Candidate skills, strict scorer, and LLM-Judge rubric; audit gate for flow-level safety.

Table 27: Experiment-level reproducibility controls. The experiments are inference and text-skill optimization runs; no hidden model parameters are trained.

Reproduction uses sanitized input tables, prediction files, policy-memory summaries, prompt templates, mediation code, and scoring code. The reported L2 Closed setting uses 40 policy rules (12 Ask, 11 Allow, and 17 Mask), local retrieval over task, app, persona, recipient, and visual evidence, and deterministic mediation. The study uses inference, replay, and text-skill optimization only; no model is trained or fine-tuned.

Component	Reproducible setting and boundary
Local VLM	8B-class local multimodal model; local inference wrapper, inference only.
Cloud VLMs	Rule-free cloud controls with raw screenshots, temperature 0.0, and parsed JSON outputs.
LLM-Judge	External multimodal judge used only for skill-evolution sandbox checks.
Rule retrieval	40-rule policy memory; local retrieval over task, app, persona, recipient, and visual evidence.
Skill evolution	Cold/medium/hot starts, up to 20 iterations, seed-controlled mutation, and validation filters.

Table 28: Model, retrieval, and implementation settings needed to interpret the reported experiments. Provider- or endpoint-specific metadata belongs in the released artifact.

L1/L3 reproducibility item	Role
L1 OCR predictions	Text-coverage denominator for the main-table L1 score.
Visible-PII L3 input	L3 denominator for visible-PII screenshots.
Sensitive items	Sensitive-item extraction output.
Strict mask audit	Reported strict SafeScreenshot audit table.
Audit manifest	Targeted visual audit pack for manual inspection.

Table 29: Key L1/L3 items for reproducing the Safe-Screenshot component audit. Exact filenames belong in the supplementary artifact.

L Dataset Card, Example Records, and Release Scope

P-GUI-Evo is a research benchmark for GUI privacy mediation, built from abstracted UI patterns, prompt-drafted scenarios, checked policy labels, sanitized sensitive fields, and reconstructed HTML/UI screenshots. Its intended use is controlled GUI privacy research; private-attribute inference and medical, financial, workplace, or identity decisions are outside scope.

If screenshots are released, they will receive a final privacy scan for accounts, credentials, QR codes, identifiers, and sensitive records; unreleased visuals will be replaced by sanitized labels, prompts, scripts, schemas, and rendering metadata where permitted.

Field group	Examples	Visibility
Runner input	screenshot, app/platform, persona context, user instruction, agent intent	method input
Policy output	Allow, Mask, Ask	prediction/evaluation audit/reproduction
Construction metadata	UI pattern, prompt template, persona/task/recipient context, HTML/UI layout, render viewport	
Sensitive-field meta-data	PII type, fake value map, injection method, localizable status	offline only
Audit metadata	risk tag, localizable region, expected label, annotation notes	offline only
Release metadata	sample id, bucket, split, sanitized scenario summary	sanitized artifact

Table 30: Dataset schema and field visibility. Deployed mediator inputs receive only runner-visible fields.

Item	Statement
Intended use	Research benchmark for GUI-agent privacy mediation, policy arbitration, SafeScreenshot construction, and skill-evolution evaluation.
Non-use boundary	Not intended for inferring real users’ private attributes, identity verification, medical or financial decision making, workplace surveillance, or training systems to expose private visual evidence.
Source and construction	Reconstructed GUI interaction patterns: LLM-assisted metadata drafting, policy-label checks, sanitized sensitive-field injection, HTML/UI rendering, and author review.
Release artifacts	Sanitized metadata, labels, prompt templates, scoring code, policy-memory summaries, and release-eligible screenshots or masked derivatives after final privacy review. During anonymous review, materials should use an anonymous supplementary artifact or placeholder.
Withheld assets	Any image that contains real accounts, contacts, organization identifiers, QR codes, credentials, identity documents, or reversible identifiers should be withheld or replaced by sanitized surrogates.
License / terms	The P-GUI-Evo benchmark release is intended under CC BY 4.0 for release-eligible reconstructed metadata, labels, templates, and sanitized visual assets. Third-party models, tools, and API services remain governed by the respective upstream licenses or provider terms, with details recorded in the artifact README/NOTICE file.

Table 31: Dataset release and intended-use statement. The release scope is deliberately narrower than the internal experimental workspace.

Example	Policy	Context	English gloss	Role
WeChat medical record transfer	Ask	User A, medical sharing to an external consultant	A reconstructed chat prepares medical details for external consultation; the system should pause and avoid silent exposure.	Ask gate / no masking example
QQ bank-card chat	Mask	User C, chat with an online seller involving a financial identifier	A reconstructed QQ conversation contains a bank-card number; the SafeScreenshot masks the financial identifier.	Mask / exposure-denial example
Hospital OA travel request	Allow	User A, structured workplace form	A reconstructed workplace form contains employee fields; the example illustrates a task-relevant UI that should remain inspectable.	Allow / no masking example

Table 32: Representative reconstructed records corresponding to Figure 6. The English gloss explains the privacy boundary while preserving the original benchmark label.






Figure 5: Qualitative SafeScreenshot examples from the strict L3 audit. The first row shows a synthetic bank-card span; the second row shows an address-list UI with multiple address and phone fields. OCR localizes matched sensitive text, and SafeScreenshot masks those spans while preserving task-relevant structure such as address-category tags. Both examples are no-flag audit cases; all shown identifiers and addresses are synthetic.

A AskConfirmation: pause for user confirmation, not redaction

Original UI




After policy gate (unchanged UI)




B Mask: SafeScreenshot comparison

Original UI




SafeScreenshot output



C Allow: task-relevant UI remains inspectable

Original UI



After policy gate (unchanged UI)




Figure 6: Representative reconstructed benchmark examples. Ask and Allow keep the original UI on both sides; Mask shows the before/after SafeScreenshot comparison. All shown visual content is synthetic or reconstructed.